Quantifying the Trendiness of Trends

Andreas Kryger Jensen and Claus Thorn Ekstrøm
Biostatistics, Institute of Public Health, University of Copenhagen

24 April, 2019

Abstract

A statement often seen in the news concerning some public health outcome is that some trend has changed or been broken. Such statements are often based on longitudinal data from e.g., surveys, and the change in the trend is claimed to have occurred at the time of the latest data collection. These types of statistical assessments are very important as they may potentially influence public health decisions on a national level.

We propose two measures for quantifying the *trendiness* of a trend. Under the assumption that reality evolves in continuous time we define what constitutes a trend and a change in a trend, and we introduce a probabilistic Trend Direction Index. This index has the intuitive interpretation of the probability that a latent characteristic has changed monotonicity at any given time conditional on observed data. We also define a global index of Expected Trend Instability quantifying the expected number of times that a trend has changed on an interval.

Using a latent Gaussian process model we show how the Trend Direction Index and the Expected Trend Instability can be estimated from data in a Bayesian framework and give an application to development of the proportion of smokers in Denmark during the last 20 years.

Keywords: Functional Data Analysis, Gaussian Processes, Trends, Bayesian Statistics, Trend Direction Index, Expected Trend Instability

1 Introduction

Trend detection has received increased attention in many fields, and while many important applications have their roots in the fields of economics (stock development) and environmental change (global temperature), trend identification has important ramifications in industry (process monitoring), medicine (disease development) and public health (changes in society).

This manuscript is concerned with the fundamental problem of estimating an underlying trend based on random variables observed repeatedly over time. In addition to this problem we also wish to assess points in time where it is possible that such a trend is changing. Our motivation is a recent example from the news in Denmark where it was stated that the trend in the proportion of smokers in Denmark has changed at the end of the year 2018 such that the proportion is now increasing whereas it had been decreasing for the previous 20 years. This statement was based on survey data collected yearly since 1998 and reported by the Danish Health Authority (The Danish Health Authority 2019), and it is critical for the Danish Health Authorities to be able to evaluate and reaqct if there is an actual change in trend (see Figure~1).

The concept of "trend" is not itself well-defined and it is often up to each researcher to define exactly what is meant by "trend" (Esterby 1993). Consequently, it is not obvious, what constitutes a *change* in trend, and that makes it difficult to compare statistical methods to detect trend changes, since they attempt to address slightly different problems.

Change point analysis is an often-used approach for detecting if a change (of a pre-specified type) has taken place (Basseville and Nikiforov 1993, @barry1993bayesianchangepoint). However, change

Figure 1: The proportion of daily or occasional smokers in Denmark during the last 20 years estimated from survey data and reported by the Danish Health Authority. [**TODO:** Mention in text why 2009 is missing]

point analysis is marred by fact that change points that happen at the boundary of the observed time frame are notoriously difficult to detect, and that they assume that change points are abrupt changes that occur at a specific time point.

Another common approach to analyse trends in time series applying a low-pass smoothing filter to the observed data in order to remove noise and extract the underlying latent trend (Chandler and Scott 2011). But in addition to the smoothing filter, it is also necessary to circumvent the problem of what consitutes a change in trend so it is necessary to specify a decision criteria to determine if any filtered changes are relevant.

Gottlieb and Müller (2012) define a stickiness coefficient for longitudinal data and use the stickiness coefficient to summarize the extent to which deviations from the mean trajectory tend to co-vary over time. While the stickness coefficient determines changes from the expected trend it is a measure that is not easily interpreted.

We propose a new method for evaluating changes in the latent trend, f, which starts by a clear specification of the problem and relevant measures of the trend: The trend direction index gives a local probability of the monotonicity of f and provides an answer to the research question „What is the probability that the latent trend is increasing?". The expected trend instability index gives the expected number of changes in the monotonicity of f on an interval and answers research questions about the rate of changes in direction of f and can be used to evaluate the statement in the new paper that „It is the first time in 20 years that f has stoped decreasing and started to increase"]. Besides providing exact answers to specific and relevant research questions, our proposed method estimates the actual probability that the trend is changing.

%Thus we are avoiding calculate the probability that an observed change in trend is %different than would have been expected if there was no change in trend.

[**TODO:** Reference some previous work]. Our ETI has some similarities to the tracking index of

Foulkes and Davis (1981).

The manuscript is structured as follows: In Section 2 we present our statistical model based on a latent Gaussian process formulation giving rise to analytic expressions for the Trend Direction Index and the Expected Trend Instability index conditional on observed data. Section 3 is concerned with estimating the models parameters, and we give an extended application to the development of the proportion of smokers in Denmark during the last 20 years in Section 4. We conclude with a discussion.

2   Methods

In the following we assume that reality evolves in continuous time $t \in \mathcal{T} \subset \mathbb{R}$ and that there exists a random, latent function $f = \{f(t) : t \in \mathcal{T}\}$ defined on a compact subset of the real line, $\mathcal{T}$, that governs the underlying evolution of some observable characteristic in a population. We can observe this latent characteristic with noise by sampling f at discrete time points according to the additive model $Y_i = f(t_i) + \varepsilon_i$ where ε_i is a zero mean random variable independent of both Y_i and $f(t_i)$. Given observations of the form $(Y_i, t_i)_{i=1}^n$ we are interested in modeling the dynamical properties of f.

The trend of f is defined as its instantaneous slope given by the function $df(t) = \left(\frac{\mathrm{d}f(s)}{\mathrm{d}s} \right)(t)$, and f is increasing and has a positive trend at t if $df(t) > 0$, and f is decreasing with a negative trend at t if $df(t) < 0$. A change in trend is defined as a change in the sign of df, i.e., when f goes from increasing to decreasing or vice versa. As f is a random function there are no exact points in time where df changes sign. There is instead a gradual and continuous change in the monotonicity of f, and an assessment of a change in trend is defined by the probability of the sign of df. This stands in contrast to traditional change-point models which assume that there are one or more exact time points where a sudden change in function or its parameterization occurs (Carlstein, Müller, and Siegmund 1994).

The probability that a trend is changing at time t is quantified by the Trend Direction Index

$$\mathrm{TDI}(t, \delta) = P(df(t + \delta) > 0 \mid \mathcal{F}_t), \quad t \in \mathcal{T} \tag{1}$$

where $\mathcal{F}_t$ is a σ-algebra of available information observed up until time t. The value of $\mathrm{TDI}(t, \delta)$ is a local probabilistic index, and it is equal to the probability that f is increasing at time $t + \delta$ given everything known about the data generating process up until and including time t. A similar definition can be given for a negative trend but that is equal to $1 - \mathrm{TDI}(t, \delta)$ and therefore redundant. The sign of δ determines whether the Trend Direction Index estimates the past ($\delta \leq 0$) or forecasts the future ($\delta > 0$). Most of the examples seen in the news concerning public health outcomes is concerned with t being equal to the current calendar time and $\delta = 0$. This excludes the usage of both change-point and segmented regression models (Quandt 1958) as there are no observations available beyond the stipulated change-point. A useful reparameterization of the Trend Direction Index is $\mathrm{TDI}(\max \mathcal{T}, t - \max \mathcal{T})$ with $t \leq \max \mathcal{T}$. This parameterization conditions on the full observation period and looks back in time whereas setting $t = \max \mathcal{T}$ corresponds to the current Trend Direction Index at the end of the observation period.

In addition to the Trend Direction Index we define a global measure of trend instability. Informally we say that a random function f is *trend stable* on an interval $\mathcal{I}$ if its sample paths maintain their monotonicity so that the trends do not change sign on the interval. To quantify the trend *in*stability

we propose to use the expected number of zero-crossings by df on $\mathcal{I}$. We define the Expected Trend Instability as

$$\text{ETI}(\mathcal{I}) = \text{E}\left[\#\left\{t \in \mathcal{I} : df(t) = 0\right\} \mid \mathcal{F}\right] \tag{2}$$

equal to the expected value of the size of the random set of zero-crossings by df on $\mathcal{I}$ when conditioning on a suitable σ-algebra $\mathcal{F}$. A common case is when $\mathcal{F}$ is generated by all observed data on $\mathcal{T}$ and $\mathcal{I} \subseteq \mathcal{T}$. The lower $\text{ETI}(\mathcal{I})$ is, the more stable is the trend of f on $\mathcal{I}$ and vice versa.

These two general definitions of trendiness will be evaluated in the light of a particular statistical model in the next sections leading to expressions of their estimates.

2.1 Latent Gaussian Process Model

The definitions in the previous section impose restrictions on the latent function f. We shall assume that f is a Gaussian process on $\mathcal{T}$. From a Bayesian perspective this is equivalent to imposing an infinite dimensional prior distribution on the latent characteristic that governs the observed outcomes. Statistical models with Gaussian process priors are a flexible approach for non-parametric regression (Radford 1999). Using a latent Gaussian process also provides an analytically tractable way for performing statistical inference on its derivatives. The general idea of our model is to apply the properties of the Gaussian process prior on on the latent characteristic to update its the finite dimensional distributions by conditioning on the observed data. This results in a posterior Gaussian process and its derivatives from which estimates of the trendiness indices will be derived.

A random function f on is a Gaussian process if and only if the vector $(f(t_1), \ldots, f(t_n))$ has a multivariate normal distribution for every finite set of evaluation points $(t_1, \ldots, t_n)$, and we write $f \sim \mathcal{GP}(\mu(\cdot), C(\cdot, \cdot))$ where μ is the mean function and C is a symmetric, positive definite covariance function (Cramer and Leadbetter 1967). We observed dependent data in terms of outcomes and their associated sampling times, $(Y_i, t_i)_{i=1}^n$, and we assume that the data are generated by the hierarchical model

$$\begin{aligned} f \mid \beta, \theta &\sim \mathcal{GP}(\mu_\beta(\cdot), C_\theta(\cdot, \cdot)) \\ Y_i \mid t_i, f(t_i), \Theta &\overset{iid}{\sim} N(f(t_i), \sigma^2), \quad \Theta = (\beta, \theta, \sigma) \end{aligned} \tag{3}$$

where β is a vector of parameters for the mean function of f, θ is a vector of parameters governing the covariance of f, and σ is the conditional standard deviation of the observations. We assume the following regularity conditions.

Assumption 1. ***TODO*** *We assume that f is a separable Gaussian process and that*

A1: ?

A2: ?

This ensures that the sample paths of f are in $C^r(\mathcal{T})$. We require that $r \geq 1$ for the Trend Direction Index and $r \geq 2$ for the index of Expected Trend Instability. [@scheuerer2010regularity]

Under the above assumptions, an important property of a Gaussian process is that it together with its first and second derivatives is a multivariate Gaussian process with explicit expressions for the joint mean, covariance and cross-covariance functions. Specifically, the joint distribution of the latent function, f, and its first and second derivatives, df and d^2f, is the multivariate Gaussian

process

$$
\begin{bmatrix} f(s) \\ df(t) \\ d^2f(u) \end{bmatrix} \mid \beta, \theta \sim \mathcal{GP} \left(\begin{bmatrix} \mu_\beta(s) \\ d\mu_\beta(t) \\ d^2\mu_\beta(u) \end{bmatrix}, \begin{bmatrix} C_\theta(s,s') & \partial_2 C_\theta(s,t) & \partial_2^2 C_\theta(s,u) \\ \partial_1 C_\theta(t,s) & \partial_1\partial_2 C_\theta(t,t') & \partial_1\partial_2^2 C_\theta(t,u) \\ \partial_1^2 C_\theta(u,s) & \partial_1^2\partial_2 C_\theta(u,t) & \partial_1^2\partial_2^2 C_\theta(u,u') \end{bmatrix} \right) \tag{4}
$$

where $d^k\mu_\beta$ is the k'th derivative of μ_β and ∂_j^k denotes the k'th order partial derivative with respect to the j'th variable (Cramer and Leadbetter 1967). Proposition 1 states the joint posterior distribution of (f, df, d^2f) conditional on the observed data.

Proposition 1. *Let the data generating model be defined as in Equation (3) and* $\mathbf{Y} = (Y_1, \ldots, Y_n)$ *the vector of observed outcomes together with its sampling times* $\mathbf{t} = (t_1, \ldots, t_n)$. *Then under [**TODO**: Assumptions ???] the joint distribution of* (f, df, d^2f) *conditional on* $\mathbf{Y}$ *and* $\mathbf{t}$ *evaluated at the vector* $\mathbf{t}^*$ *of p time points is*

$$
\begin{bmatrix} f(\mathbf{t}^*) \\ df(\mathbf{t}^*) \\ d^2f(\mathbf{t}^*) \end{bmatrix} \mid \mathbf{Y}, \mathbf{t}, \Theta \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma})
$$

where $\boldsymbol{\mu} \in \mathbb{R}^{3p}$ *is the column vector of posterior expectations and* $\boldsymbol{\Sigma} \in \mathbb{R}^{3p \times 3p}$ *is the joint posterior covariance matrix. Partitioning these as*

$$
\boldsymbol{\mu} = \begin{bmatrix} \mu_f(\mathbf{t}^* \mid \Theta) \\ \mu_{df}(\mathbf{t}^* \mid \Theta) \\ \mu_{d^2f}(\mathbf{t}^* \mid \Theta) \end{bmatrix}, \quad \boldsymbol{\Sigma} = \begin{bmatrix} \Sigma_f(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) & \Sigma_{f,df}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) & \Sigma_{f,d^2f}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) \\ \Sigma_{f,df}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta)^T & \Sigma_{df}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) & \Sigma_{df,d^2f}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) \\ \Sigma_{f,d^2f}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta)^T & \Sigma_{df,d^2f}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta)^T & \Sigma_{d^2f}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) \end{bmatrix}
$$

the individual components are given by

$$
\mu_f(\mathbf{t}^* \mid \Theta) = \mu_\beta(\mathbf{t}^*) + C_\theta(\mathbf{t}^*, \mathbf{t}) \left(C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I \right)^{-1} (\mathbf{Y} - \mu_\beta(\mathbf{t}))
$$

$$
\mu_{df}(\mathbf{t}^* \mid \Theta) = d\mu_\beta(\mathbf{t}^*) + \partial_1 C_\theta(\mathbf{t}^*, \mathbf{t}) \left(C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I \right)^{-1} (\mathbf{Y} - \mu_\beta(\mathbf{t}))
$$

$$
\mu_{d^2f}(\mathbf{t}^* \mid \Theta) = d^2\mu_\beta(\mathbf{t}^*) + \partial_1^2 C_\theta(\mathbf{t}^*, \mathbf{t}) \left(C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I \right)^{-1} (\mathbf{Y} - \mu_\beta(\mathbf{t}))
$$

$$
\Sigma_f(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) = C_\theta(\mathbf{t}^*, \mathbf{t}^*) - C_\theta(\mathbf{t}^*, \mathbf{t}) \left(C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I \right)^{-1} C_\theta(\mathbf{t}, \mathbf{t}^*)
$$

$$
\Sigma_{df}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) = \partial_1\partial_2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) - \partial_1 C_\theta(\mathbf{t}^*, \mathbf{t}) \left(C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I \right)^{-1} \partial_2 C_\theta(\mathbf{t}, \mathbf{t}^*)
$$

$$
\Sigma_{d^2f}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) = \partial_1^2\partial_2^2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) - \partial_1^2 C_\theta(\mathbf{t}^*, \mathbf{t}) \left(C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I \right)^{-1} \partial_2^2 C_\theta(\mathbf{t}, \mathbf{t}^*)
$$

$$
\Sigma_{f,df}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) = \partial_2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) - C_\theta(\mathbf{t}^*, \mathbf{t}) \left(C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I \right)^{-1} \partial_2 C_\theta(\mathbf{t}, \mathbf{t}^*)
$$

$$
\Sigma_{f,d^2f}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) = \partial_2^2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) - C_\theta(\mathbf{t}^*, \mathbf{t}) \left(C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I \right)^{-1} \partial_2^2 C_\theta(\mathbf{t}, \mathbf{t}^*)
$$

$$
\Sigma_{df,d^2f}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) = \partial_1\partial_2^2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) - \partial_1 C_\theta(\mathbf{t}^*, \mathbf{t}) \left(C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I \right)^{-1} \partial_2^2 C_\theta(\mathbf{t}, \mathbf{t}^*)
$$

Proof. *See Appendix .1.* ∎

Proposition 1 grants the foundation for the rest of the methodological development. In the following two subsection we show how the Trend Direction Index and the Expected Trend Instability can be expressed under the data generating model in Equation (3) using the results of Proposition 1.

2.2 The Trend Direction Index

The Trend Direction Index was defined generally in Equation (1). Conditioning on the σ-algebra of all observed data we may express the Trend Direction Index under the model in Equation (3) through the posterior distribution of df. The following proposition states this result.

Proposition 2. *Let $\mathcal{F}_t$ be the σ-algebra generated by $(\mathbf{Y}, \mathbf{t})$ where $\mathbf{Y} = (Y_1, \ldots, Y_n)$ is the vector of observed outcomes and $\mathbf{t} = (t_1, \ldots, t_n)$ the associated sampling times. [**TODO**: Assumptions] The Trend Direction Index defined in Equation (1) can then be written in terms of the posterior distribution of df as*

$$\mathrm{TDI}(t, \delta \mid \Theta) = P(df(t + \delta) > 0 \mid \mathbf{Y}, \mathbf{t}, \Theta)$$
$$= \frac{1}{2} + \frac{1}{2} \mathrm{Erf}\left(\frac{\mu_{df}(t + \delta \mid \Theta)}{2^{1/2} \Sigma_{df}(t + \delta, t + \delta \mid \Theta)^{1/2}} \right)$$

where $\mathrm{Erf} \colon x \mapsto 2\pi^{-1/2} \int_0^x \exp(-u^2) \mathrm{d}u$ is the error function and μ_{df} and Σ_{df} are the posterior mean and covariance of the trend defined in Proposition 1.

Proof. *The proof follows directly from the definition of the Trend Direction Index and integration of a Gaussian density function.* ∎

Proposition 2 shows that the Trend Direction Index is equal to 0.5 when $\mu_{df}(t + \delta \mid \Theta) = 0$ corresponding to when the expected value of the posterior of f is locally constant at time $t + \delta$. A decision rule based on $\mathrm{TDI}(t, \delta \mid \Theta) \lessgtr 50\%$ is therefore a natural choice for assessing the local trendiness of f. However, different thresholds based on external loss or utility functions can be used depending on the application.

The Trend Direction Index is illustrated in Figure 2 in the noise free case, $\sigma = 0$, with known values of β and θ, and where the prior expected value of f is set equal to $\mu_\beta(t) = 0$. The first and second rows of the plot show 150 random sample paths from the posterior distribution of (f, df) with the posterior expectations in bold lines, and the third row shows the Trend Direction Index. Since the parameters, Θ, of f are known in this example, the Trend Direction Index is a deterministic function of time. The three columns in the plot show how the posterior of f, df, and the TDI are updated after one, two and four observations both forwards and backwards in time. The figure shows how the inclusion of additional observations results in changes of the posterior distribution of the trend and the Trend Direction Index. The variation of the forecasts remains stable, whereas the posterior distribution back in time restricts the variation of the curves to accommodate the observations. The vertical dotted lines denote the point in time after which forecasting occurs. When forecasting beyond the last observation, the posterior of f becomes more and more dominated by its prior implying that df becomes symmetric around zero so that the Trend Direction Index stabilizes around 50%.

2.3 The Expected Trend Instability

The Expected Trend Instability was defined in Equation (2) as the expected number of roots of the trend on an interval conditional on observed data. Now we make this concept more precise and frame it in the context of the latent Gaussian process model. Let $\mathcal{I}$ be a compact subset of the real line and consider the random càdlàg function

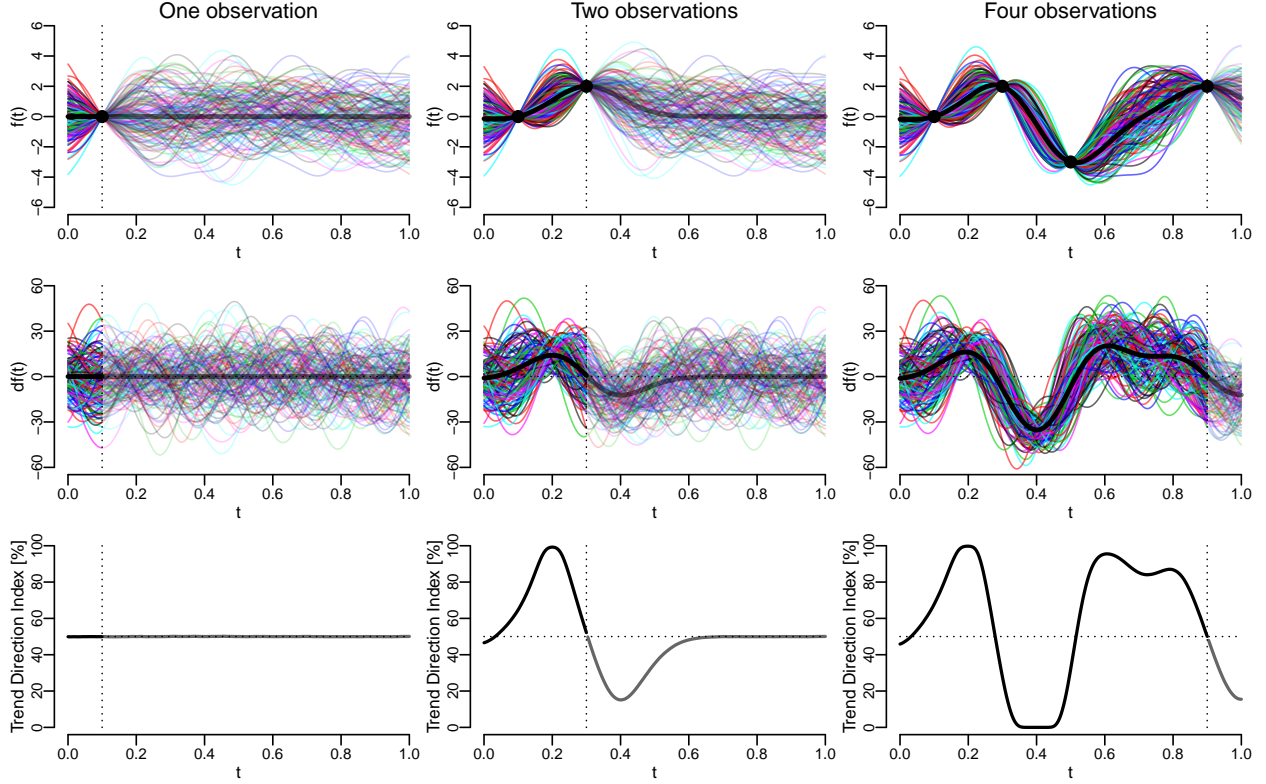$$N_\mathcal{I}(t) = \#\{s \leq t : df(s) = 0, s \in \mathcal{I}\}$$

Figure 2: 150 realizations from the posterior distribution of f (top row), df (middle row) with expected values in bold and the Trend Direction Index (bottom row) conditional on one, two and four noise free observations. Dotted vertical lines show the points in time after which forecasting takes place.

counting the cumulative number of points in $\mathcal{I}$ up to time t where the trend is equal to zero. [**TODO**: Something about tangents and way it's the number of crossings] The Expected Trend Instability on $\mathcal{I}$ is equal to

$$\mathrm{ETI}(\mathcal{I} \mid \Theta) = \mathrm{E}\left[N_{\mathcal{I}}\left(\max \mathcal{I}\right) \mid \Theta, \mathcal{F}\right]$$

giving the expected total number of zero-crossings by df on $\mathcal{I}$. The following proposition gives the expression for the Expected Trend Instability under the data generating model in Equation (3).

Proposition 3. *Let $\mathcal{F}$ be the σ-algebra generated by the observed data $(\mathbf{Y}, \mathbf{t})$ and μ_{df}, μ_{d^2f}, Σ_{df}, Σ_{d^2f} and Σ_{df,d^2f} the moments of the joint posterior distribution of (df, d^2f) as stated in Proposition 1. [**TODO**: Assumptions.] The Expected Trend Instability is then*

$$\mathrm{ETI}(\mathcal{I} \mid \Theta) = \int_{\mathcal{I}} d\mathrm{ETI}(t \mid \Theta) \mathrm{d}t$$

where $d\mathrm{ETI}$ is the local Expected Trend Instability given by

$$d\mathrm{ETI}(t, \mathcal{T} \mid \Theta) = \lambda(t \mid \Theta)\phi\left(\frac{\mu_{df}(t \mid \Theta)}{\Sigma_{df}(t, t \mid \Theta)^{1/2}}\right)\left(2\phi(\zeta(t \mid \Theta)) + \zeta(t \mid \Theta)\,\mathrm{Erf}\left(\frac{\zeta(t \mid \Theta)}{2^{1/2}}\right)\right)$$

and $\phi\colon x \mapsto 2^{-1/2}\pi^{-1/2}\exp(-\frac{1}{2}x^2)$ is the standard normal density function, $\mathrm{Erf}\colon x \mapsto$

7

$2\pi^{-1/2}\int_0^x \exp(-u^2)\mathrm{d}u$ *is the error function, and* λ, ω *and* ζ *are functions defined as*

$$\lambda(t \mid \Theta) = \frac{\Sigma_{d^2f}(t,t \mid \Theta)^{1/2}}{\Sigma_{df}(t,t \mid \Theta)^{1/2}} \left(1 - \omega(t \mid \Theta)^2\right)^{1/2}$$

$$\omega(t \mid \Theta) = \frac{\Sigma_{df,d^2f}(t,t \mid \Theta)}{\Sigma_{df}(t,t \mid \Theta)^{1/2}\Sigma_{d^2f}(t,t \mid \Theta)^{1/2}}$$

$$\zeta(t \mid \Theta) = \frac{\mu_{df}(t \mid \Theta)\Sigma_{d^2f}(t,t \mid \Theta)^{1/2}\omega(t)\Sigma_{df}(t,t \mid \Theta)^{-1/2} - \mu_{d^2f}(t \mid \Theta)}{\Sigma_{d^2f}(t,t \mid \Theta)^{1/2} \left(1 - \omega(t \mid \Theta)^2\right)^{1/2}}$$

*The value of ETI is a finite [**TODO:** finite because of assumptions] real positive value.*
Proof. *See Appendix .2.* ∎

To illustrate the Expected Trend Instability, Figure 3 shows 25 random Gaussian processes on the unit interval simulated under three different values of ETI. The sample paths are paired so that each function in the first row has an associated trend function in the second row. The different values of ETI are set to 0.25, 0.5 and 1 corresponding to the expected number of times that the functions change monotonicity or equivalently that the trend crosses zero on that interval. Sample paths that are *trend stable*, i.e., always increasing/decreasing, are shown by solid blue lines, and sample paths that are *trend unstable*, i.e., the derivatives crosses zero, are shown by dashed gold colored lines. It is seen that for low values of ETI most of the sample paths preserve their monotonicity on the interval and their associated derivatives are correspondingly either always positive or negative. For larger values of ETI, more of the trends start crossing zero implying less stability in the monotonicity of f. We note that even though we are only modeling a single curve, the Expected Trend Instability is defined in terms of a posterior distribution of random curves which is what the figure illustrates.

3 Estimation

The trendiness indices depend on the parameters Θ of the latent Gaussian process. These parameters must be estimated from the observed data, and we consider two different estimation procedures: maximum likelihood (also called Empirical Bayes) estimation and a fully Bayesian estimator.

The maximum likelihood estimator consists of finding the values of the parameters that maximize the marginal likelihood of the observed data and plugging these into the expressions of the posterior distributions and the trendiness indices. The marginal distribution of $\mathbf{Y}$ is found by integrating out the distribution of the latent Gaussian process in the conditional observation model $\mathbf{Y} \mid f(\mathbf{t}), \mathbf{t}, \Theta$ in Equation (3). Since the observation model consists of normal distributed random variables conditional on the latent Gaussian process, the marginal distribution is multivariate normal with expectation $\mu_\beta(\mathbf{t})$ and $n \times n$ covariance matrix $C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I$. The marginal log likelihood function is therefore

$$\log L(\Theta \mid \mathbf{Y}, \mathbf{t}) \propto -\frac{1}{2}\log|C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I| - \frac{1}{2}(\mathbf{Y} - \mu_\beta(\mathbf{t}))^T \left[C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I\right]^{-1}(\mathbf{Y} - \mu_\beta(\mathbf{t})) \quad (5)$$

and the maximum likelihood estimate $\widehat{\Theta^{\mathrm{ML}}} = \arg\sup_{\Theta=(\beta,\theta,\sigma)} \log L(\Theta \mid \mathbf{Y}, \mathbf{t})$ can be obtained by numerical optimization or found as the roots to the score equations $\nabla_\theta \log L(\Theta \mid \mathbf{Y}, \mathbf{t}) = 0$. This estimate can then be plugged in to the expressions for the posterior distributions of (f, df, d^2f) in Proposition 1 enabling simulation of the posterior distributions at any vector of time points. Point estimates of the Trend Direction Index and the Expected Trend Instability are then $\mathrm{TDI}(t, \delta \mid \widehat{\Theta^{\mathrm{ML}}})$ and $\mathrm{ETI}(\mathcal{I} \mid \widehat{\Theta^{\mathrm{ML}}})$ according to Propositions 2 and 3 respectively.

Figure 3: 25 random pairs sampled from the joint distribution of a Gaussian Process (f) and its derivative (df) with different values of Expected Trend Instability (ETI). The first row shows samples from f, and the second row shows samples from df. The columns define different values of ETI. Sample paths that are trend stable are shown by solid blue lines, and unstable sample paths are shown by dashed gold colored lines.

The maximum likelihood estimator is easy to implement and fast to compute, but it is difficult to propagate the uncertainties of the parameter estimates through to the posterior distributions and the trendiness indices. This is disadvantageous since in order to conduct a qualified assessment of trendiness we are not only interested in point estimates but also the associated uncertainties. A Bayesian estimator, while slower to compute, facilitates a straightforward way to encompass all uncertainties in the final estimates. To specify a Bayesian estimator we must augment the data generating model in Equation (3) with another hierarchical level specifying a prior distribution of the model parameters. We therefore add the following level

$$(\beta, \theta, \sigma) \sim G(\Theta \mid \Psi, \mathbf{t})$$

to the specification where G is some family of distribution functions indexed by a vector of hyper-parameters Ψ. The joint distribution of the model can be factorized as

$$P(\mathbf{Y}, f(\mathbf{t}), \Theta \mid \Psi, \mathbf{t}) = P(\mathbf{Y} \mid f(\mathbf{t}), \Theta, \Psi, \mathbf{t})P(f(\mathbf{t}) \mid \Theta, \Psi, \mathbf{t})G(\Theta \mid \Psi, \mathbf{t})$$

and each conditional probability is specified by a sub-model in the augmented hierarchy. We always condition on Ψ and $\mathbf{t}$ as they are considered fixed data. The posterior distribution of Θ given the observed outcomes is then

$$
\begin{aligned}
P(\Theta \mid \mathbf{Y}, \Psi, \mathbf{t}) &= \frac{G(\Theta \mid \Psi, \mathbf{t})P(\mathbf{Y} \mid \Theta, \Psi, \mathbf{t})}{P(\mathbf{Y} \mid \Psi, \mathbf{t})} \\
&= \frac{G(\Theta \mid \Psi, \mathbf{t}) \int P(\mathbf{Y} \mid f(\mathbf{t}), \Theta, \Psi, \mathbf{t})dP(f(\mathbf{t}) \mid \Theta, \Psi, \mathbf{t})}{\iint P(\mathbf{Y} \mid f(\mathbf{t}), \Theta, \Psi, \mathbf{t})dP(f(\mathbf{t}) \mid \Theta, \Psi, \mathbf{t})dG(\Theta \mid \Psi, \mathbf{t})}
\end{aligned}
$$

9

Let $\widetilde{\Theta} \sim P(\Theta \mid \mathbf{Y}, \Psi, \mathbf{t})$. The posterior distribution of Θ induces corresponding distributions over the trendiness indices according to $\mathrm{TDI}(t, \delta \mid \widetilde{\Theta})$, $d\mathrm{EDI}(t, \mathcal{T}, \mid \widetilde{\Theta})$ and $\mathrm{EDI}(\mathcal{I} \mid \widetilde{\Theta})$. For example, the Trend Direction Index in the Bayesian formulation is a surface in (t, δ) where each value is a distribution over probability values. We suggest to summarize the trendiness indices by their posterior quantiles. For example, for the Trend Direction Index we summarize its posterior distribution by functions Q_τ such that

$$P\left(\mathrm{TDI}\left(t, \delta \mid \widetilde{\Theta}\right) \leq \tau\right) = Q_\tau(t, \delta)$$

with for example $\tau \in \{0.025, 0.5, 0.975\}$ to obtain 95% credibility intervals.

We have implemented both the maximum likelihood and the Bayesian estimator in Stan (Carpenter et al. 2017) and R (R Core Team 2018) in combination with the `rstan` package (Stan Development Team 2018). Stan is a probabilistic programming language enabling full Bayesian inference using Markov chain Monte Carlo sampling. The Stan implementation of the maximum likelihood estimator requires the marginal maximum likelihood estimates of the parameters supplied as data, and from these it will simulate random functions from the posterior distribution of $(f, df, d^2 f)$ on a user-supplied grid of time points as well as return point estimates of TDI and dETI. The latter can then be integrated numerically to obtain the Expected Trend Instability on an interval. The Bayesian estimator requires values of the hyper-parameters Ψ supplied as data and from these it will generate random samples $(\widetilde{\Theta}_1, \ldots, \widetilde{\Theta}_K)$ from $P(\Theta \mid \mathbf{Y}, \Psi, \mathbf{t})$ by Markov chain Monte Carlo. These samples are then used to approximate the posterior distribution of the Trend Direction Index by $(\mathrm{TDI}(t, \delta \mid \widetilde{\Theta}_1), \ldots, \mathrm{TDI}(t, \delta \mid \widetilde{\Theta}_K))$ and similarly for dETI.

4 Application

A report published by The Danish Health Authority in January 2019 updated the estimated proportion of daily or occasional smokers in Denmark with new data from 2018 (The Danish Health Authority 2019). The data was based on an online survey including 5017 participants. The report also included data on the proportion of smokers in Denmark during the last 20 years which was shown in Figure 1. The report was picked up by several news papers under headlines stating that the proportion of smokers in Denmark had significantly increased for the first time in two decades (Navne, Schmidt, and Rasmussen n.d.). The report published no statistical analyses for this statement but wrote, that because the study population is so large, then more or less all differences become statistically significant at the 5% level (this was written as a 95% significance level in the report).

This data set provides an instrumental way of exemplifying our two proposed measures of trendiness. In this application we wish to assess the statistical properties of following questions:

Q1: Is the proportion of smokers increasing in the year 2018 conditional on data from the last 20 years?

Q2: If the proportion of smokers is currently increasing, and if yes: when did this increase probably start?

Q3: Is it the first time during the last 20 years that the trend in the proportion of smokers has changed?

A naive approach for trying to answer questions Q1 and Q2 is to apply sequential χ^2-tests in a 2×2 table. Table 1 shows the p-values for the χ^2-test (with Yates' continuity correction) of independence

between the proportion of smokers in 2018 and each of the five previous years. Using a significance level of 5% the conclusion is ambiguous. Compared to the previous year, there was no significant change in the proportion in 2018. Three out of these five comparisons fail to show a significant change in proportions. It is therefore evident that such point-wise testing is not sufficiently perceptive to catch the underlying continuous development.

2018	2017	2016	2015	2014	2013
p-value	0.074	**0.020**	0.495	**0.012**	0.576

Table 1: p-values obtained from χ^2-tests of independence between the proportion of smokers in 2018 and the five previous years. Numbers in bold are statistically significant differences at the 5% level.

Similarly, a simple approach for trying to answer question Q3 would be to look at the cumulative number of times that the difference in proportion between consecutive years changes sign. In the data set there were nine changes in the sign of the difference between the proportion in each year and the proportion in the previous year giving this very crude estimate of the number of times that the trend has changed. This approach suffers from the facts that it is based on a finite difference approximation at the sampling points to the continuous derivative, and that it uses the noisy measurements instead the latent function. Consequently, the changes in trend can be quite unstable.

We now present an analysis of the data set using our method. To complete the model specification in Equation (3) we must decide on the functional forms of the mean and covariance functions of the latent Gaussian process. We considered three different mean functions: a constant mean: $\mu_\beta(t) = \beta_0$, a linear mean: $\mu_\beta(t) = \beta_0 + \beta_1 t$, and a quadratic mean: $\mu_\beta(t) = \beta_0 + \beta_1 t + \beta_2 t^2$. We considered the Squared Exponential (SE), the Rational Quadratic (RQ), the Matern 3/2 (M3/2), and the Matern 5/2 (M5/2) covariance functions (Rasmussen and Williams 2006). These covariance functions are

$$
\begin{aligned}
C_\theta^{\text{SE}}(s,t) &= \alpha^2 \exp\left(-\frac{(s-t)^2}{2\rho^2}\right), \quad \theta = (\alpha, \rho) > 0 \\
C_\theta^{\text{RQ}}(s,t) &= \alpha^2 \left(1 + \frac{(s-t)^2}{2\rho^2 \nu}\right)^{-\nu}, \quad \theta = (\alpha, \rho, \nu) > 0 \\
C_\theta^{\text{M3/2}}(s,t) &= \alpha^2 \left(1 + \frac{\sqrt{3}\sqrt{(s-t)^2}}{\rho}\right) \exp\left(-\frac{\sqrt{3}\sqrt{(s-t)^2}}{\rho}\right), \quad \theta = (\alpha, \rho) > 0 \\
C_\theta^{\text{M5/2}}(s,t) &= \alpha^2 \left(1 + \frac{\sqrt{5}\sqrt{(s-t)^2}}{\rho} + \frac{5(s-t)^2}{3\rho^2}\right) \exp\left(-\frac{\sqrt{5}\sqrt{(s-t)^2}}{\rho}\right), \quad \theta = (\alpha, \rho) > 0
\end{aligned}
\tag{6}
$$

We note that the Rational Quadratic covariance function converges on $\mathcal{T} \times \mathcal{T}$ to the Squared Exponential covariance function for $\nu \to \infty$. [**TODO:** Something about the sample path regularity for these covariance functions. M3/2 cannot be used for ETI but okay for TDI].

The total combination of mean and covariance functions results in 12 candidate models, and to select the model providing the best fit to data we used leave-one-out cross-validation based on maximum likelihood estimation. For each model, $\mathcal{M}$, we turn by turn excluded data from a single year, and we let the leave-one-out data be denoted $(\mathbf{Y}_{-i}, \mathbf{t}_{-i})$. Based on the leave-one-out data sets we estimated the parameters $\Theta_{-i}^{\mathcal{M}} = (\beta, \theta, \sigma^2)$ for each model by maximizing the marginal log-likelihood. These estimates are given by

$$
\widehat{\Theta}_{-i}^{\mathcal{M}} = \arg\sup_{\Theta} \log L(\Theta \mid \mathbf{Y}_{-i}, \mathbf{t}_{-i})
$$

where the log likelihood function is given in Equation (5). These estimates were then plugged into the expression for the posterior expectation of f in Proposition 1 to obtain the leave-one-out predictions at times t_i. The mean squared error of prediction was calculated as the squared errors between the leave-one-out predictions and the observed values averaged across all data points

$$\text{MSPE}_{\text{LOO}}^{\mathcal{M}} = \frac{1}{n} \sum_{i=1}^{n} \left(Y_i - \text{E}[f(t_i) \mid \mathbf{Y}_{-i}, \mathbf{t}_{-i}, \widehat{\Theta}_{-i}^{\mathcal{M}}] \right)^2$$

and the model providing the best fit to data was selected as $\mathcal{M}_{\text{opt}} = \arg\min_{\mathcal{M}} \text{MSPE}_{\text{LOO}}^{\mathcal{M}}$.

	SE	RQ	Matern 3/2	Matern 5/2
Constant	0.682	**0.651**	0.687	0.660
Linear	0.806	$\Leftarrow$	0.896	0.865
Quadratic	0.736	$\Leftarrow$	0.800	0.785

Table 2: Leave-one-out cross-validated mean squared error of prediction for each of the 12 candidate models. $\Leftarrow$ indicates numerical convergence to the SE covariance function.

Table 2 shows the mean squared error of prediction for each candidate model. For the models with a Rational Quadratic covariance function and a linear and a quadratic mean function the parameter ν diverged numerically implying convergence to the Squared Exponential covariance function. Comparing the leave-one-out mean squared error of prediction, the prior distribution of f in the optimal model has a constant mean function and a Rational Quadratic covariance function. The marginal maximum likelihood estimates of the parameters in the optimal model were

$$\widehat{\beta_0^{\text{ML}}} = 28.001, \quad \widehat{\alpha^{\text{ML}}} = 4.543, \quad \widehat{\rho^{\text{ML}}} = 4.438, \quad \widehat{\nu^{\text{ML}}} = 1.020, \quad \widehat{\sigma^{\text{ML}}} = 0.622 \tag{7}$$

Figure 4 shows the fit of the model by the maximum likelihood method. The plots were obtained by plugging the maximum likelihood estimates into the expressions for the posterior distributions of f and df defined in Proposition 1, the Trend Direction Index in Proposition 2, and the Expected Trend Instability in Proposition 3. The predictions were performed on an equidistant grid of 500 time points spanning the 20 years. The plot of the posterior trend (top right) shows two regions in time where the posterior mean of the derivative is positive, one around $[2004; 2008]$ and one shortly after 2015 and until the end of the observation period. The Trend Direction Index (bottom left) quantifies this positive trendiness as a probability standing in 2018 and looking back in time while also taking the uncertainty into account. The bottom right panel shows the local Expected Trend Instability and its integral is to the expected number of times that the trend has changed sign. Table 3 summarizes the maximum likelihood estimates of the Trend Direction Index at the end of the observation period and the previous five years as well as the Expected Trend Instability during the full observation period and during only the last ten years. Crosspoint is the first point in time during the last ten years where the Trend Direction Index became greater than 50% i.e.,

$$\text{Crosspoint} = \arg\min_{t \in [-10; 0]} \left\{ 2018 + t : \text{TDI}(2018, 2018 - t) \geq 50\% \right\}$$

Based on the results from the maximum likelihood analysis we may answer the questions by stating that the proportion of smokers in Denmark is currently increasing with a probability of 95.24%. This is, however, not a recent development as the probability of an increasing proportion has been greater than 50% since the middle of 2015. This can be compared to the sequential χ^2-tests in

12

Table 1, which gave a cruder and less consistant result. The estimated values of ETI in Table 3 show that there has been an average of 3.68 changes in the monotonicity of the proportion during the last 20 years. This value does support the statement by the news outlets that it is the first time in 20 years that the trend has changed. The value, however, reduces to 1.39 when only looking 10 years back, which is slightly less than half the ETI for the longer period.
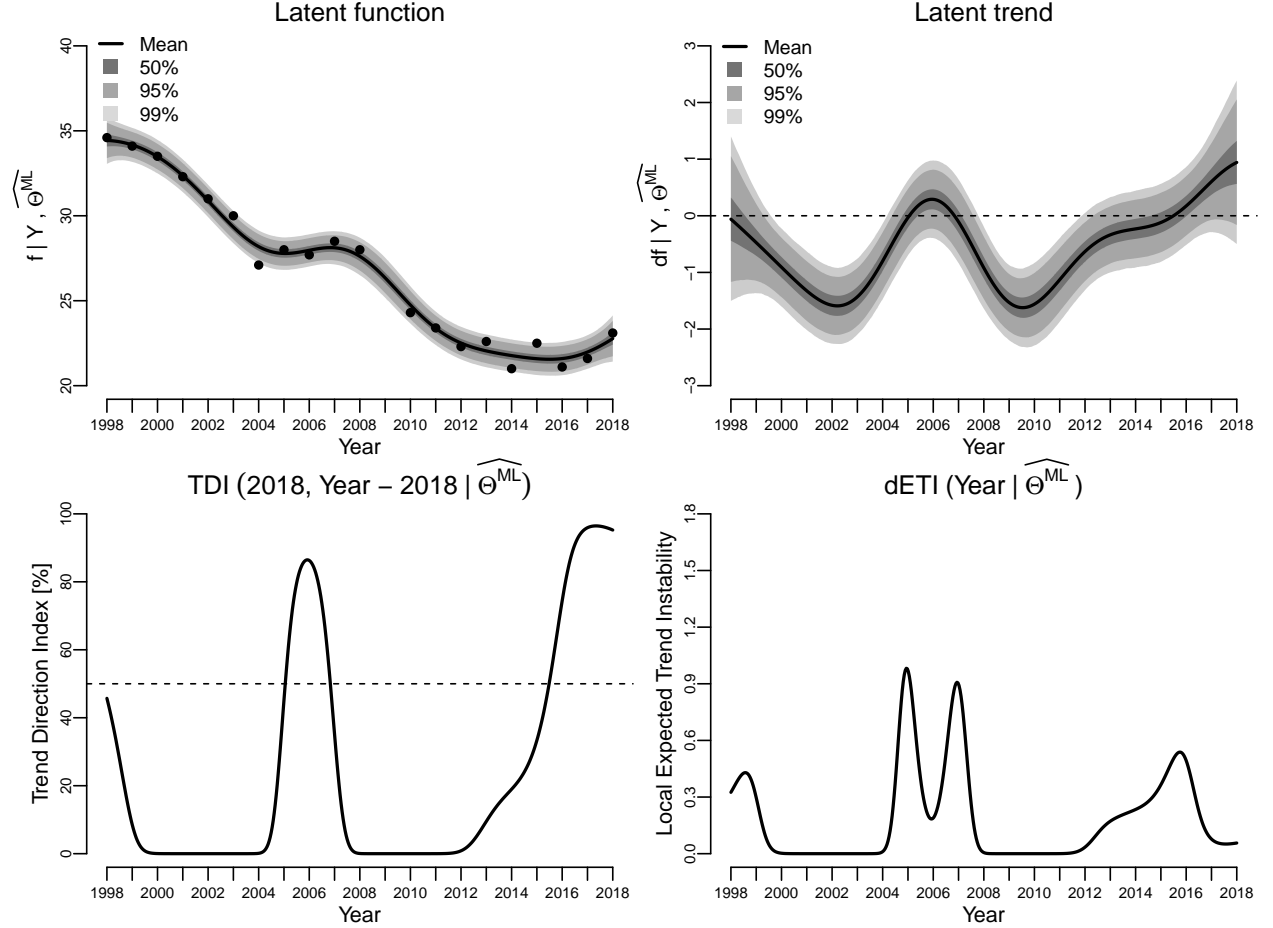

Figure 4: Results from fitting the latent Gaussian process model by maximum likelihood. The first row shows the posterior distributions of f (left) and df (right) with the posterior means in bold and gray areas showing point-wise probability intervals for the posterior distribution. The second row shows the estimated Trend Direction Index (left) and the local Expected Trend Instability (right).

We also applied the Bayesian estimator to the data using the same prior mean and covariance structure. The Bayesian estimator requires a prior distribution for the model parameters. We used independent priors of the form

$$G(\Theta \mid \Psi, \mathbf{t}) = G(\beta_0 \mid \Psi_{\beta_0}, \mathbf{t})G(\alpha \mid \Psi_\alpha, \mathbf{t})G(\rho \mid \Psi_\rho, \mathbf{t})G(\nu \mid \Psi_\nu, \mathbf{t})G(\sigma \mid \Psi_\sigma, \mathbf{t})$$

where each prior was a heavy-tailed distribution with a moderate variance centered at the maximum likelihood estimates. We used the following distributions

$$\beta_0 \sim T\left(\widehat{\beta_0^{\mathrm{ML}}}, 3, 3\right), \qquad \alpha \sim \text{Half-}T\left(\widehat{\alpha^{\mathrm{ML}}}, 3, 3\right), \quad \rho \sim \text{Half-}N\left(\widehat{\rho^{\mathrm{ML}}}, 1\right)$$

$$\nu \sim \text{Half-}T\left(\widehat{\nu^{\mathrm{ML}}}, 3, 3\right), \quad \sigma \sim \text{Half-}T\left(\widehat{\sigma^{\mathrm{ML}}}, 3, 3\right)$$

where the maximum likelihood values are given in Equation (7) and Half-$T(\cdot, \cdot, \mathrm{df})$ and Half-$N(\cdot, \cdot)$ denotes the location-scale half T- and normal distribution functions with df degrees of freedom due to the requirement of positivity. We used a more informative prior for ρ compared to the other parameters, since it was less identified from data. We ran four independent Markov chains for 25,000 iterations each with half of the iterations used for warm-up and discarded. Convergence was assessed by the potential scale reduction factor, $\widehat{R}$, of Gelman and Rubin (1992).

Figure 4 shows results from the Bayesian estimator. In this model both trendiness indices are time-varying posterior distributions and the top row of the shows the posterior distributions of the Trend Direction Index (left) and the Local Expected Trend Instability (right) summarized by time-dependent quantiles. The bottom row shows posterior density estimates of the Expected Trend Instability during the last twenty years (left) and during the last ten years (right). The same summary statistics as for the maximum likelihood analysis are given in Table 3 but here stated in terms of posterior medians and 2.5% and 97.5% posterior quantiles.
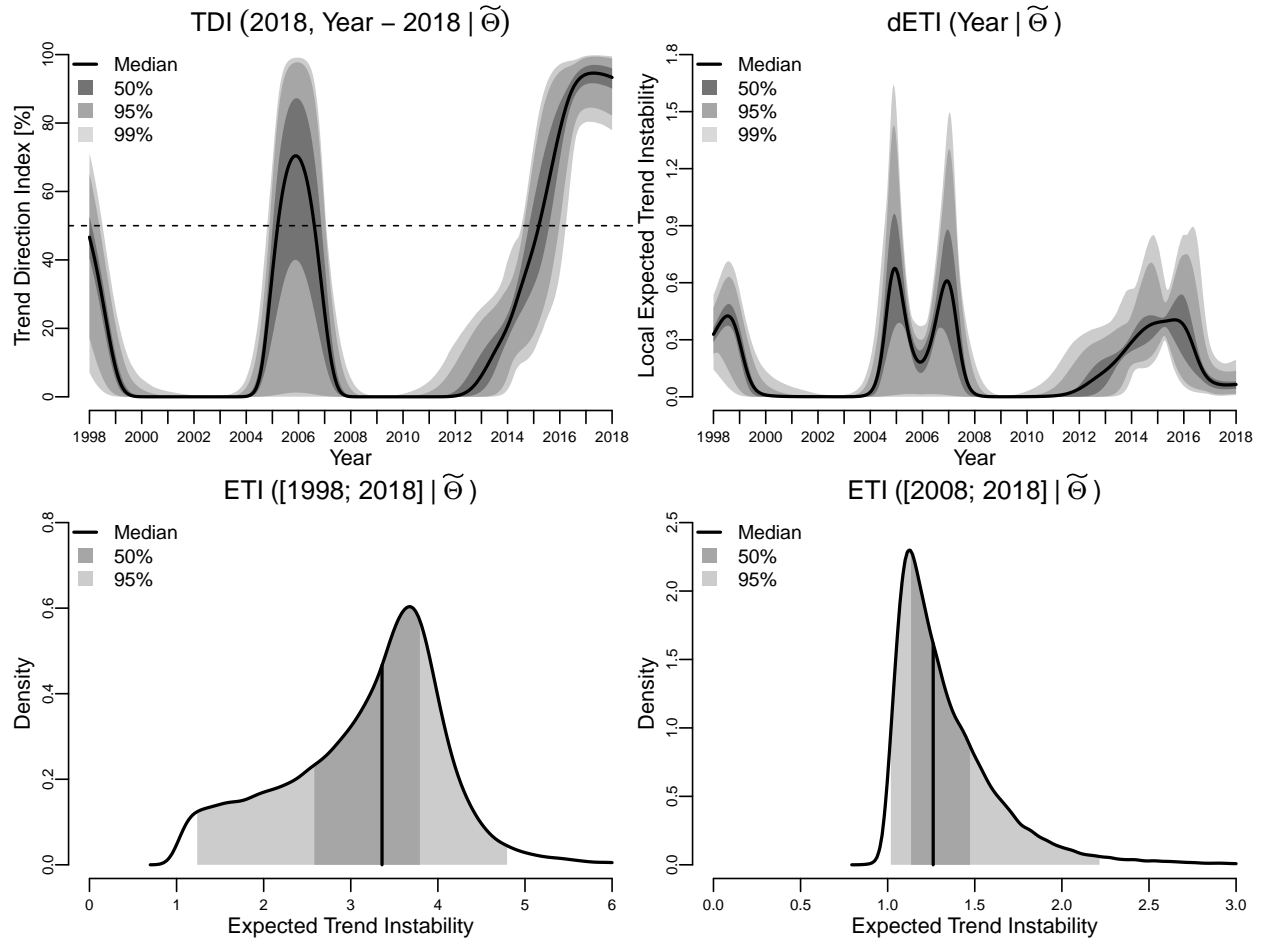
\begin{figure}[htb]

\caption{Results from fitting the latent Gaussian process model by Bayesian analysis. The first row shows the posterior distributions of TDI (left) and local ETI (right) with the posterior means in bold and gray areas showing point-wise 95% and 99% probability intervals for the posterior distribution. The second row shows densities and probability intervals for the expected trend instability for the 20-year period back-in-time from 2018 (left) and 10 year back-in-time (right).}
\end{figure}

	Maximum Likelihood	Bayesian Posterior	
TDI$(2018, 0)$	95.24%	93.32%	$[82.15\%; 98.86\%]$
TDI$(2018, -1)$	95.92%	94.21%	$[84.28\%; 99.11\%]$
TDI$(2018, -2)$	74.41%	77.87%	$[51.02\%; 94.94\%]$
TDI$(2018, -3)$	33.36%	44.11%	$[18.23\%; 69.19\%]$
TDI$(2018, -4)$	18.96%	20.60%	$[6.05\%; 31.82\%]$
TDI$(2018, -5)$	9.50%	6.21%	$[0.03\%; 22.21\%]$
Crosspoint	2015.48	2015.19	$[2014.62; 2015.96]$
ETI$([1998, 2018])$	3.68	3.36	$[1.24; 4.79]$
ETI$([2008, 2018])$	1.39	1.25	$[1.02; 2.22]$

Table 3: Summary measures from the Maximum Likelihood and Bayesian analyses. The rows show the estimated Trend Direction Index for 2013 to 2018 and the Expected Trend Instability for the last 10 and 20 years all conditional on data from 1998 to 2018. For the Bayesian analysis posterior medians and 95% posterior probability intervals are given. Crosspoint is the time during the last 10 years where TDI first exceeded 50%.

The results from the two analyses generally agree, but there are two differences that we wish to address. Both analyses showed a local peak in trendiness around 2015. In the maximum likelihood analysis this occurred at 2005.94 with a Trend Direction Index of 86.47%. In the Bayesian analysis the peak occurs at 2005.87 with a median Trend Direction Index of 79.43% and a 95% posterior probability interval of $[1.23\%; 97.73\%]$. The added uncertainty estimates facilitated by the Bayesian estimator shows this trendiness is so variable that there is no reason to believe in a substantial increase in proportions at that point in time. This insight could not have been obtained from the maximum likelihood analysis.

The second difference is that the Bayesian model seems to generally induce a high degree of smoothness in the estimates [**TODO**: smoothness is the wrong word here]. This can be seen from the plot of the median local Expected Trend Instability in Figure 4 which is generally lower than its corresponding maximum likelihood point estimates in Figure 4. This is similarly reflected in the median ETI estimates in Table 3 which are lower than their values under maximum likelihood. Looking at the posterior distributions of the covariance parameters θ (not shown), we see that this is mainly a result of not restricting the parameter ν to its maximum likelihood value. The 95% probability interval of the posterior distribution of ν was $[0.328; 10.743]$ which is highly right-skewed compared to the maximum likelihood estimate of $\widehat{\nu^{\mathrm{ML}}} = 1.020$.

To understand the effect of ν on the smoothness of the fitted models we can compare the local expected number of crossings by a Gaussian process and its derivative at their mean values in the simple case of a zero-mean process with either the Rational Quadratic or the Squared Exponential covariance function. In this case the formula in Proposition 3 simplifies immensely, and as shown in Appendix .3 the local expected number of mean-crossings by f is equal to $\pi \rho^{-1}$ for both covariance functions. However, for df the local expected number of mean-crossings is equal to $3^{1/2}\pi^{-1}\rho^{-1}$ for the Squared Exponential covariance function and $3^{1/2}\pi^{-1}\rho^{-1}(1 + \nu^{-1})^{1/2}$ for the Rational Quadratic covariance function. We note that $1 < (1 + v^{-1})^{1/2} < \infty$ for $0 < \nu < \infty$ and monotonically decreasing for $\nu \to \infty$ with a limit of one. Therefore, the value of ν has no effect on the crossing intensity of the process itself, but its derivative is always larger under a Rational Quadratic covariance function compared to the Squared Exponential covariance function with equality in the limit. A right-skewed posterior distribution of ν therefore favors fewer crossings of

15

the trend leading to a more stable trend and a smaller value of the Expected Trend Instability.

5 Discussion

[**TODO:** Her må der båttes løs]

We are computing a different (but more interpretable) probability than what can be obtained from e.g., a χ^2 test. So it's fundamentally different. Also, we are able to exploit the underlying assumed continuity.

Something about that the indices are scale-free hence TDI should be interpreted along with $df \mid Y$ at the same time point. We could also define

$$\text{TDI}_u(t, \delta) = P(df(t + \delta) > u \mid \mathcal{F}_t)$$

where u is some trend value.

[**TODO:** Assert that both TDI and EDI are scale free. They don't tell anything about the magnitude of a trend. They should therefore always be accompanied by plots of the posterior of df].

We could also have defined

$$\text{ETI}^{\uparrow}(\mathcal{I}) = \text{E}\left[\# \left\{t \in \mathcal{I} : df(t) = 0, d^2f(t) > 0\right\} \mid \mathcal{F}\right]$$

counting the expected number of up-crossings at zero by df (the number of times f has gone from being decreasing to increasing) on $\mathcal{T}$.

Reproducible code and Stan implementations are available at Jensen (n.d.).

#Bibliography

Barry, Daniel, and J. A. Hartigan. 1993. "A Bayesian Analysis for Change Point Problems." *Journal of the American Statistical Association* 88 (421): 309–19. http://www.jstor.org/stable/2290726.

Basseville, Michele, and Igor V. Nikiforov. 1993. *Detection of Abrupt Changes: Theory and Application.* Prentice-Hall.

Carlstein, Edward, Hans-Georg Müller, and David Siegmund, eds. 1994. *Change-Point Problems.* Vol. 23. Lecture Notes – Monograph Series. Institute of Mathematical Statistics.

Carpenter, Bob, Andrew Gelman, Matthew D Hoffman, Daniel Lee, Ben Goodrich, Michael Betancourt, Marcus Brubaker, Jiqiang Guo, Peter Li, and Allen Riddell. 2017. "Stan: A Probabilistic Programming Language." *Journal of Statistical Software* 76 (1).

Chandler, R., and M. Scott. 2011. *Statistical Methods for Trend Detection and Analysis in the Environmental Sciences.* Statistics in Practice. Wiley. https://books.google.dj/books?id=zLh3OkKWJ7QC.

Cramer, Harald, and M. R. Leadbetter. 1967. *Stationary and Related Stochastic Processes – Sample Function Properties and Their Applications.* John Wiley & Sons, Inc.

Esterby, S. R. 1993. "Trend Analysis Methods for Environmental Data." *Environmetrics* 4 (4): 459–81. https://doi.org/10.1002/env.3170040407.

Foulkes, Mary Ann, and CE Davis. 1981. "An Index of Tracking for Longitudinal Data." *Biometrics*, 439–46.

Gelman, Andrew, and Donald B. Rubin. 1992. "Inference from Iterative Simulation Using Multiple Sequences." *Statistical Science* 7 (4): 457–72.

Gottlieb, Andrea, and Hans-Georg Müller. 2012. "A Stickiness Coefficient for Longitudinal Data." *Computational Statistics & Data Analysis* 56 (12): 4000–4010.

Gradshteyn, Izrail Solomonovich, and Iosif Moiseevich Ryzhik. 2014. *Table of Integrals, Series, and Products*. Academic Press.

Jensen, Andreas Kryger. n.d. "GitHub Repository for the Trendiness of Trends." Accessed 2019. https://github.com/aejensen/TrendinessOfTrends.

Navne, Helene, Anders Legarth Schmidt, and Lars Igum Rasmussen. n.d. "Første Gang I 20 år: Flere Danskere Ryger." Accessed 2019. https://politiken.dk/forbrugogliv/sundhedogmotion/art6938627/Flere-danskere-ryger.

Quandt, Richard E. 1958. "The Estimation of the Parameters of a Linear Regression System Obeying Two Separate Regimes." *Journal of the American Statistical Association* 53 (284): 873–80.

Radford, Neal M. 1999. "Regression and Classification Using Gaussian Process Priors (with Discussion)." In *Bayesian Statistics 6: Proceedings of the Sixth Valencia International Meeting*, edited by A. P. Dawid José M. Bernardo James O. Berger and Adrian F. M. Smith, 475–501.

Rasmussen, C. E., and C. K. I. Williams. 2006. *Gaussian Processes in Machine Learning*. MIT Press.

R Core Team. 2018. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. https://www.R-project.org/.

Rice, Stephen O. 1945. "Mathematical Analysis of Random Noise, II." *Bell System Technical Journal* 24 (1): 46–156.

Stan Development Team. 2018. "RStan: The R Interface to Stan." http://mc-stan.org/.

The Danish Health Authority. 2019. "Danskernes Rygevaner 2018." 2019. www.sst.dk/da/udgivelser/2019/danskernes-rygevaner-2018.

#Appendix

.1   Proof of Proposition 1

Let $\mathbf{Y} = (Y_1, \ldots, Y_n)$ and $\mathbf{t} = (t_1, \ldots, t_n)$ be the vectors of observed outcomes and associated sampling times. From the data generating model in Equation (3) we observe that the marginal distribution of the vector of observed outcomes $\mathbf{Y} \mid \mathbf{t}, \Theta$ is

$$P(\mathbf{Y} \mid \mathbf{t}, \Theta) = \int P(\mathbf{Y} \mid f(\mathbf{t}), \mathbf{t}, \Theta) dP(f(\mathbf{t}) \mid \mathbf{t}, \Theta)$$
$$= N(\mu_\beta(\mathbf{t}), C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I)$$

where $\mu_\beta(\mathbf{t}) = (\mu_\beta(t_1), \ldots \mu_\beta(t_n))$, $C_\theta(\mathbf{t}, \mathbf{t})$ is the $n \times n$ covariance matrix obtained by evaluating $C_\theta(s, t)$ at $\{(s, t) \in \mathbf{t} \times \mathbf{t}\}$ and I is an $n \times n$ identity matrix. This implies that the joint distribution

of $\mathbf{Y}$ and the latent functions (f, df, d^2f) evaluated at an arbitrary vector of time points $\mathbf{t}^*$ is

$$
\begin{bmatrix} f(\mathbf{t}^*) \\ df(\mathbf{t}^*) \\ d^2f(\mathbf{t}^*) \\ \mathbf{Y} \end{bmatrix} \mid \mathbf{t}, \Theta \sim N \left(\begin{bmatrix} \mu_\beta(\mathbf{t}^*) \\ d\mu_\beta(\mathbf{t}^*) \\ d^2\mu_\beta(\mathbf{t}^*) \\ \mu_\beta(\mathbf{t}) \end{bmatrix}, \begin{bmatrix} C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_2^2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) & C_\theta(\mathbf{t}^*, \mathbf{t}) \\ \partial_1 C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_1\partial_2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_1\partial_2^2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_1 C_\theta(\mathbf{t}^*, \mathbf{t}) \\ \partial_1^2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_1^2\partial_2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_1^2\partial_2^2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_1^2 C_\theta(\mathbf{t}^*, \mathbf{t}) \\ C_\theta(\mathbf{t}, \mathbf{t}^*) & \partial_2 C_\theta(\mathbf{t}, \mathbf{t}^*) & \partial_2^2 C_\theta(\mathbf{t}, \mathbf{t}^*) & C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I \end{bmatrix} \right)
$$

where ∂_j^k denotes the k'th order partial derivative with respect to the j'th variable.

By the standard formula for deriving conditional distributions in a multivariate normal model, the posterior distribution of (f, df, d^2f) evaluated at the p time points in $\mathbf{t}^*$ is

$$
\begin{bmatrix} f(\mathbf{t}^*) \\ df(\mathbf{t}^*) \\ d^2f(\mathbf{t}^*) \end{bmatrix} \mid \mathbf{Y}, \mathbf{t}, \Theta \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma})
$$

where $\boldsymbol{\mu} \in \mathbb{R}^{3p}$ is the column vector of posterior expectations and $\boldsymbol{\Sigma} \in \mathbb{R}^{3p \times 3p}$ is the joint posterior covariance matrix, and these are given by

$$
\boldsymbol{\mu} = \begin{bmatrix} \mu_\beta(\mathbf{t}^*) \\ d\mu_\beta(\mathbf{t}^*) \\ d^2\mu_\beta(\mathbf{t}^*) \end{bmatrix} + \begin{bmatrix} C_\theta(\mathbf{t}^*, \mathbf{t}) \\ \partial_1 C_\theta(\mathbf{t}^*, \mathbf{t}) \\ \partial_1^2 C_\theta(\mathbf{t}^*, \mathbf{t}) \end{bmatrix} K_{\theta,\sigma}(\mathbf{t}, \mathbf{t})^{-1}(\mathbf{Y} - \mu_\beta(\mathbf{t}))
$$

$$
\boldsymbol{\Sigma} = \begin{bmatrix} C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_2^2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) \\ \partial_1 C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_1\partial_2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_1\partial_2^2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) \\ \partial_1^2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_1^2\partial_2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) & \partial_1^2\partial_2^2 C_\theta(\mathbf{t}^*, \mathbf{t}^*) \end{bmatrix} - \begin{bmatrix} C_\theta(\mathbf{t}^*, \mathbf{t}) \\ \partial_1 C_\theta(\mathbf{t}^*, \mathbf{t}) \\ \partial_1^2 C_\theta(\mathbf{t}^*, \mathbf{t}) \end{bmatrix} K_{\theta,\sigma}(\mathbf{t}, \mathbf{t})^{-1} \begin{bmatrix} C_\theta(\mathbf{t}, \mathbf{t}^*) \\ \partial_2 C_\theta(\mathbf{t}, \mathbf{t}^*) \\ \partial_2^2 C_\theta(\mathbf{t}, \mathbf{t}^*) \end{bmatrix}^T
$$

where $K_{\theta,\sigma}(\mathbf{t}, \mathbf{t}) = C_\theta(\mathbf{t}, \mathbf{t}) + \sigma^2 I$. Partitioning $\boldsymbol{\mu}$ and $\boldsymbol{\Sigma}$ as

$$
\boldsymbol{\mu} = \begin{bmatrix} \mu_f(\mathbf{t}^* \mid \Theta) \\ \mu_{df}(\mathbf{t}^* \mid \Theta) \\ \mu_{d^2f}(\mathbf{t}^* \mid \Theta) \end{bmatrix}, \quad \boldsymbol{\Sigma} = \begin{bmatrix} \Sigma_f(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) & \Sigma_{f,df}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) & \Sigma_{f,d^2f}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) \\ \Sigma_{f,df}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) & \Sigma_{df}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) & \Sigma_{df,d^2f}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) \\ \Sigma_{d^2f,f}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) & \Sigma_{d^2f,df}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) & \Sigma_{d^2f}(\mathbf{t}^*, \mathbf{t}^* \mid \Theta) \end{bmatrix}
$$

and completing the matrix algebra, we obtain the expressions of the individual components given in Proposition 1. This completes the proof.

.2 Proof of Proposition 3

Rice showed in section 3.3. of Rice (1945) that the expected number of zero-crossings of a Gaussian process X on an interval $\mathcal{I}$ is given by

$$
\int_{\mathcal{I}} \int_{-\infty}^{\infty} |v| f_{X(t),dX(t)}(0, v) \mathrm{d}v \mathrm{d}t \tag{8}
$$

where $f_{X(t),dX(t)}$ is the joint density function of X and its derivative dX at time t. [**TODO:** Conditions] To derive the expression for the Expected Trend Instability in Equation (**??**) we must apply the Rice formula to the joint posterior distribution of (df, d^2f). From Proposition 1 the distribution of $(df, d^2f) \mid \mathbf{Y}, \mathbf{t}, \Theta$ is bivariate normal for each t.

Let μ_{df}, μ_{d^2f}, Σ_{df} and Σ_{d^2f} be defined as in Proposition 1 and define further

$$
\omega(t \mid \Theta) = \frac{\Sigma_{df,d^2f}(t, t \mid \Theta)}{\Sigma_{df}(t, t \mid \Theta)^{1/2} \Sigma_{d^2f}(t, t \mid \Theta)^{1/2}}
$$

as the posterior point-wise cross-correlation function between df and d^2f. The joint posterior density function of (df, d^2f) at any time t evaluated at $(0, v)$ can be factorized as

$$f_{df(t),d^2f(t)}(0,v) = c_1(t)e^{c_2(t)}e^{-c_3(t)v^2 - 2c_4(t)v}$$

where $c_1, \ldots, c_4$ are functions of time given by

$$c_1(t) = (2\pi)^{-1}\Sigma_{df}(t,t \mid \Theta)^{-1/2}\Sigma_{d^2f}(t,t \mid \Theta)^{-1/2}(1 - \omega(t \mid \Theta)^2)^{-1/2}$$

$$c_2(t) = \frac{\mu_{df}(t \mid \Theta)^2}{2\Sigma_{df}(t,t \mid \Theta)(\omega(t \mid \Theta)^2 - 1)} + \frac{\mu_{d^2f}(t \mid \Theta)^2}{2\Sigma_{d^2f}(t,t \mid \Theta)(\omega(t \mid \Theta)^2 - 1)}$$

$$- \frac{\mu_{df}(t \mid \Theta)\mu_{d^2f}(t \mid \Theta)\omega(t \mid \Theta)}{\Sigma_{df}(t,t \mid \Theta)^{1/2}\Sigma_{d^2f}(t,t \mid \Theta)^{1/2}(\omega(t \mid \Theta)^2 - 1)}$$

$$c_3(t) = -\frac{1}{2}\Sigma_{d^2f}(t,t \mid \Theta)^{-1}(\omega(t \mid \Theta)^2 - 1)^{-1}$$

$$c_4(t) = -\frac{\mu_{df}(t \mid \Theta)\Sigma_{d^2f}(t,t \mid \Theta)^{1/2}\omega(t \mid \Theta) - \mu_{d^2f}(t \mid \Theta)\Sigma_{df}(t,t \mid \Theta)^{1/2}}{2\Sigma_{d^2f}(t,t \mid \Theta)(\omega(t \mid \Theta)^2 - 1)\Sigma_{df}(t,t \mid \Theta)^{1/2}}$$

Let $d\mathrm{ETI}(t \mid \Theta)$ denote the inner integral in Equation (8). Using the factorization of the joint posterior density we may write it was

$$d\mathrm{ETI}(t \mid \Theta) = \int_{-\infty}^{\infty} |v| f_{df(t),d^2f(t)}(0,v)\mathrm{d}v$$

$$= c_1(t)e^{c_2(t)} \int_{-\infty}^{\infty} |v|e^{-c_3(t)v^2 - 2c_4(t)v}\mathrm{d}v \tag{9}$$

$$= c_1(t)e^{c_2(t)} \left(\int_0^{\infty} ve^{-c_3(t)v^2 + 2c_4(t)v}\mathrm{d}v + \int_0^{\infty} ve^{-c_3(t)v^2 - 2c_4(t)v}\mathrm{d}v \right)$$

Because $c_3(t) > 0$ for all t since $\Sigma_{d^2f}(t,t \mid \Theta) > 0$ and $|\omega(t \mid \Theta)| < 1$ by Assumption [**TODO**:???] we obtain the following solution for the type of integral in the previous display by using formula 5 in section 3.462 on page 365 of Gradshteyn and Ryzhik (2014)

$$\int_0^{\infty} ve^{-c_3(t)v^2 \pm 2c_4(t)v}\mathrm{d}v = \frac{1}{2c_3(t)} \pm \frac{c_4(t)}{2c_3(t)}\frac{\pi^{1/2}}{c_3(t)^{1/2}}e^{\frac{c_4(t)^2}{c_3(t)}}\left(1 \pm \mathrm{Erf}\left(\frac{c_4(t)}{\sqrt{c_3(t)}}\right)\right) \tag{10}$$

where $\mathrm{Erf}\colon x \mapsto 2\pi^{-1}\int_0^x e^{-u^2}\mathrm{d}u$ is the error function. Combining Equations (9) and (10) we may express $d\mathrm{ETI}$ as

$$d\mathrm{ETI}(t \mid \Theta) = c_1(t)e^{c_2(t)}\left(\frac{1}{c_3(t)} + \frac{c_4(t)}{c_3(t)}\frac{\pi^{1/2}}{c_3(t)^{1/2}}e^{\frac{c_4(t)^2}{c_3(t)}}\mathrm{Erf}\left(\frac{c_4(t)}{\sqrt{c_3(t)}}\right)\right)$$

Defining $\zeta(t \mid \Theta) = \sqrt{2}c_4(t)c_3(t)^{-1/2}$ and collecting some terms, the index can be rewritten as

$$d\mathrm{ETI}(t \mid \Theta) = \frac{c_1(t)}{c_3(t)}\left(e^{c_2(t)} + \frac{\pi^{1/2}}{2^{1/2}}e^{\frac{c_4(t)^2}{c_3(t)} + c_2(t)}\zeta(t)\,\mathrm{Erf}\left(\frac{\zeta(t \mid \Theta)}{2^{1/2}}\right)\right)$$

Straightforward arithmetic calculations show that

$$\frac{c_4(t)^2}{c_3(t)} + c_2(t) = -\frac{\mu_{df}(t \mid \Theta)^2}{2\Sigma_{df}(t,t \mid \Theta)}, \quad c_2(t) = -\frac{1}{2}\left(\zeta(t \mid \Theta)^2 + \frac{\mu_{df}(t \mid \Theta)^2}{\Sigma_{df}(t,t \mid \Theta)}\right)$$

and by defining $\phi\colon x \mapsto (2\pi)^{-1/2}e^{-x^2}$ as the density function of the standard normal distribution we may write $e^{\frac{c_4(t)^2}{c_3(t)}+c_2(t)} = (2\pi)^{1/2}\phi\left(\frac{\mu_{df}(t|\Theta)}{\Sigma_{df}(t,t|\Theta)^{1/2}}\right)$ and $e^{c_2(t)} = 2\pi\phi(\zeta(t))\phi\left(\frac{\mu_{df}(t|\Theta)}{\Sigma_{df}(t,t|\Theta)^{1/2}}\right)$ which leads to

$$d\mathrm{ETI}(t \mid \Theta) = \frac{c_1(t)}{c_3(t)}\pi\phi\left(\frac{\mu_{df}(t \mid \Theta)}{\Sigma_{df}(t,t \mid \Theta)^{1/2}}\right)\left(2\phi(\zeta(t \mid \Theta)) + \zeta(t \mid \Theta)\,\mathrm{Erf}\left(\frac{\zeta(t \mid \Theta)}{2^{1/2}}\right)\right)$$

Standard arithmetics show that

$$\frac{c_1(t)}{c_3(t)} = \frac{1}{\pi}\frac{\Sigma_{d^2f}(t,t \mid \Theta)^{1/2}}{\Sigma_{df}(t,t \mid \Theta)^{1/2}}\left(1 - \omega(t \mid \Theta)^2\right)^{1/2}$$

and we finally obtain the expression

$$d\mathrm{ETI}(t \mid \Theta) = \lambda(t \mid \Theta)\phi\left(\frac{\mu_{df}(t \mid \Theta)}{\Sigma_{df}(t,t \mid \Theta)^{1/2}}\right)\left(2\phi(\zeta(t \mid \Theta)) + \zeta(t \mid \Theta)\,\mathrm{Erf}\left(\frac{\zeta(t \mid \Theta)}{2^{1/2}}\right)\right)$$

where λ and ζ are given by

$$\lambda(t \mid \Theta) = \frac{\Sigma_{d^2f}(t,t \mid \Theta)^{1/2}}{\Sigma_{df}(t,t \mid \Theta)^{1/2}}\left(1 - \omega(t \mid \Theta)^2\right)^{1/2}$$

$$\zeta(t \mid \Theta) = \frac{\mu_{df}(t \mid \Theta)\Sigma_{d^2f}(t,t \mid \Theta)^{1/2}\omega(t)\Sigma_{df}(t,t \mid \Theta)^{-1/2} - \mu_{d^2f}(t \mid \Theta)}{\Sigma_{d^2f}(t,t \mid \Theta)^{1/2}\left(1 - \omega(t \mid \Theta)^2\right)^{1/2}}$$

By definition

$$\mathrm{ETI}(\mathcal{I} \mid \Theta) = \int_{\mathcal{I}} d\mathrm{ETI}(t \mid \Theta)\mathrm{d}t$$

which completes the proof.

.3 Zero-crossings of f and df in the zero-mean stationary case

Let $f \sim \mathcal{GP}\left(0, C_\theta(\cdot,\cdot)\right)$ where the C_θ is either the Squared Exponential or Rational Quadratic covariance function is given in Equation 6. We look at the expected number of zero-crossings on an interval by either f and df as given by the Rice formula in Equation (8) with either $X(t) = f(t)$ or $X(t) = df(t)$. In this case the expressions simplifies immensely due to the zero means of both f, df, and d^2f and because $\mathrm{Cov}[f(t), df(t)] = 0$ and $\mathrm{Cov}[df(t), d^2f(t)] = 0$. The latter is a result of using a stationary covariance function for the prior distribution of f (Cramer and Leadbetter 1967). In this stationary case local expected number of zero-crossing of f and df are given by

$$\frac{\left.\partial_1\partial_2 C_\theta(s,t)\right|_{s=t}^{1/2}}{\pi C_\theta(t,t)^{1/2}} \quad \text{and} \quad \frac{\left.\partial_1^2\partial_2^2 C_\theta(s,t)\right|_{s=t}^{1/2}}{\pi\left.\partial_1\partial_2 C_\theta(s,t)\right|_{s=t}^{1/2}}$$

respectively. From Equation (6) it follows that

$$C_\theta^{\mathrm{SE}}(t,t) = \sigma^2, \quad \left.\partial_1\partial_2 C_\theta^{\mathrm{SE}}(s,t)\right|_{s=t} = \frac{\sigma^2}{\rho^2}, \quad \left.\partial_1^2\partial_2^2 C_\theta^{\mathrm{SE}}(s,t)\right|_{s=t} = \frac{3\sigma^2}{\rho^4}$$

$$C_\theta^{\mathrm{RQ}}(t,t) = \sigma^2, \quad \left.\partial_1\partial_2 C_\theta^{\mathrm{RQ}}(s,t)\right|_{s=t} = \frac{\sigma^2}{\rho^2}, \quad \left.\partial_1^2\partial_2^2 C_\theta^{\mathrm{RQ}}(s,t)\right|_{s=t} = \frac{2\sigma^2(1+\nu)}{\nu\rho^4}$$

and the local expected number of zero-crossings of f and df for either the Squared Exponential and the Rational Quadratic covariance functions are

$$\frac{\partial_1\partial_2 C_\theta^{\mathrm{SE}}(s,t)\Big|_{s=t}^{1/2}}{\pi C_\theta^{\mathrm{SE}}(t,t)^{1/2}} = \frac{1}{\pi\rho}, \qquad \frac{\partial_1^2\partial_2^2 C_\theta^{\mathrm{SE}}(s,t)\Big|_{s=t}^{1/2}}{\pi\partial_1\partial_2 C_\theta^{\mathrm{SE}}(s,t)\Big|_{s=t}^{1/2}} = \frac{3^{1/2}}{\pi\rho}$$

$$\frac{\partial_1\partial_2 C_\theta^{\mathrm{RQ}}(s,t)\Big|_{s=t}^{1/2}}{\pi C_\theta^{\mathrm{RQ}}(t,t)^{1/2}} = \frac{1}{\pi\rho}, \qquad \frac{\partial_1^2\partial_2^2 C_\theta^{\mathrm{RQ}}(s,t)\Big|_{s=t}^{1/2}}{\pi\partial_1\partial_2 C_\theta^{\mathrm{RQ}}(s,t)\Big|_{s=t}^{1/2}} = \frac{3^{1/2}}{\pi\rho}\left(1+v^{-1}\right)^{1/2}$$